UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

3253 Analytic Techniques and Machine Learning

Module 6: Support Vector Machines

Course Plan

Module Titles
Module 1 – Introduction to Machine Learning
Module 2 – End to End Machine Learning Project
Module 3 – Classification
Module 4 – Clustering and Unsupervised Learning
Module 5 – Training Models and Feature Selection
Current Focus: Module 6 – Support Vector Machines
Module 7 – Decision Trees and Ensemble Learning
Module 8 – Dimensionality Reduction
Module 9 – Introduction to TensorFlow
Module 10 – Introduction to Deep Learning and Deep Neural Networks
Module 11 – Distributing TensorFlow, CNNs and RNNs
Module 12 – Final Assignment and Presentations (no content)

Learning Outcomes for this Module

- Apply linear and non-linear Support Vector Machine (SVM) to classification problems
- Use SVM for regression problems
- Explore how SVM works

Topics for this Module

- **6.1** Introduction to Support Vector Machines (SVM)
- **6.2** Linear classification with SVM
- **6.3** Non-linear classification
- **6.4** SVM regression
- **6.5** How SVM works
- **6.6** Resources and Wrap-up

Module 6 – Section 1

Introduction to Support Vector Machines (SVM)

SVM

- An ML algorithm for either
  - Classification
  - Regression
- It can model both
  - Linear
  - Non-linear patterns
- Most appropriate for small or mid-size data sets
- Training points (x, y) have labels +1 or -1

UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 6 – Section 2

Linear Classification with SVM

Linear Classification

- A training set is linearly separable if a hyperplane leaves class +1 on one side and class -1 on the other side
- SVM finds the hyperplane that leaves the largest possible margin on both sides, until the first examples are found

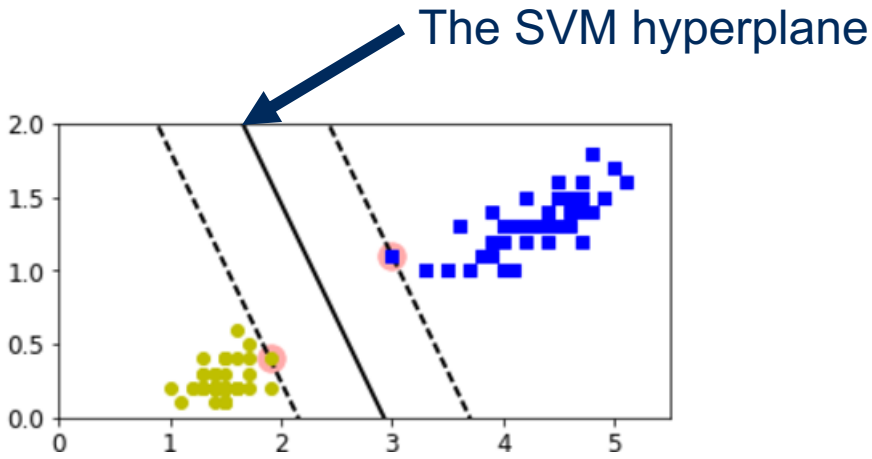

Linear Classification

- A training set is linearly separable if a hyperplane leaves class +1 on one side and class -1 on the other side
- SVM finds the hyperplane that leaves the largest possible margin on both sides, until the first examples are found

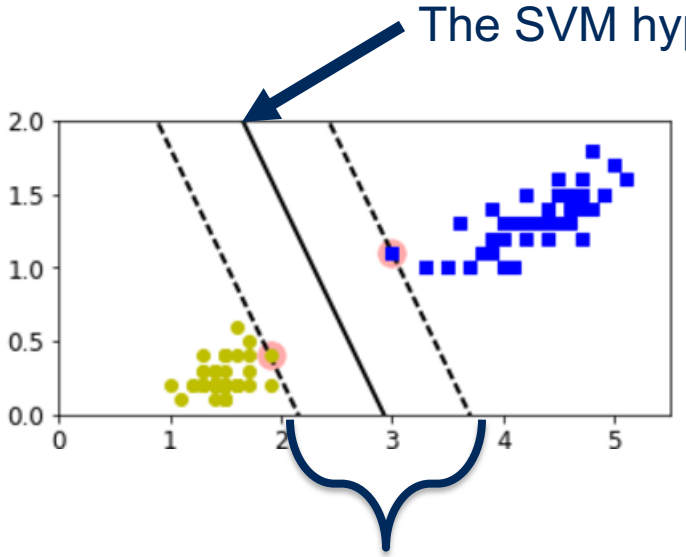

SVM aims to maximize the street width

Linear Classification

- A training set is linearly separable if a hyperplane leaves class +1 on one side and class -1 on the other side
- SVM finds the hyperplane that leaves the largest possible margin on both sides, until the first examples are found

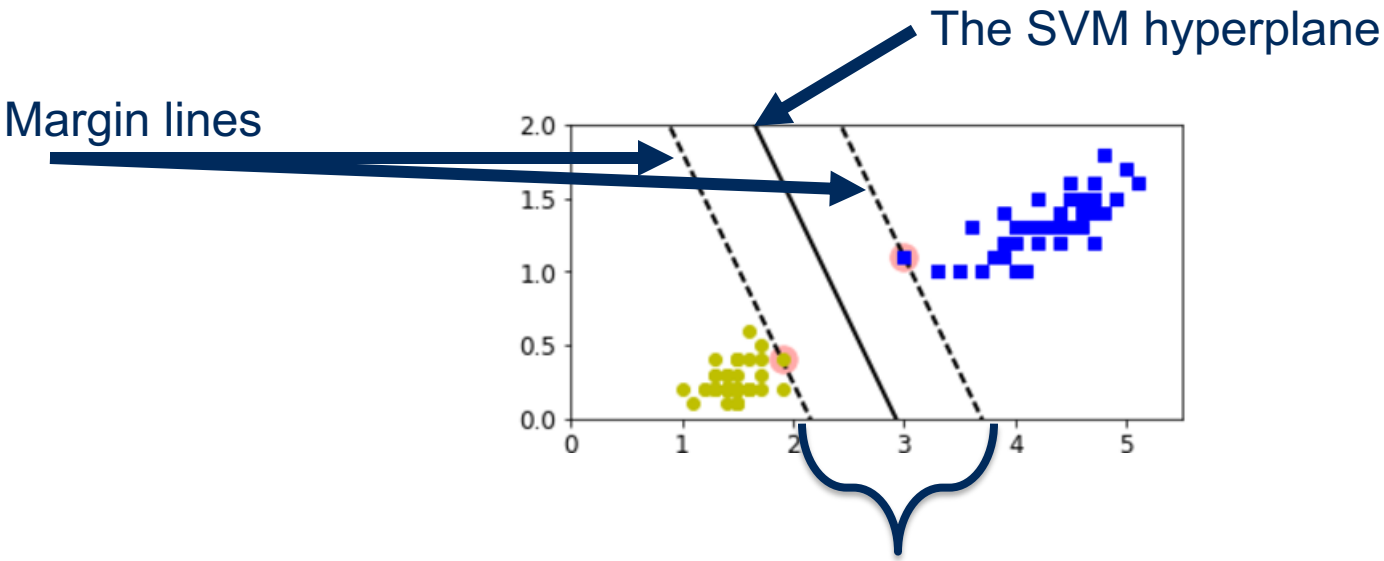

SVM aims to maximize the street width

Linear Classification (cont'd)

- A hyperplane in the plane is a line
- The points (x_1, x_2) in the line satisfy the equation

$$w_1 x_1 + w_2 x_2 + b = 0$$, w1,w2, b constants

- (w_1, w_2) is perpendicular ($\perp$) to the SVM [solid] line
- A hyperplane in 3D space is a plane described as

$$w_1 x_1 + w_2 x_2 + w_3 x_3 + b = 0$$

- In n-dimensions (features), a hyperplane is

$$w_1 x_1 + \ldots + w_n x_n + b = w^T \cdot x + b = 0$$

Linear Classification (cont'd)

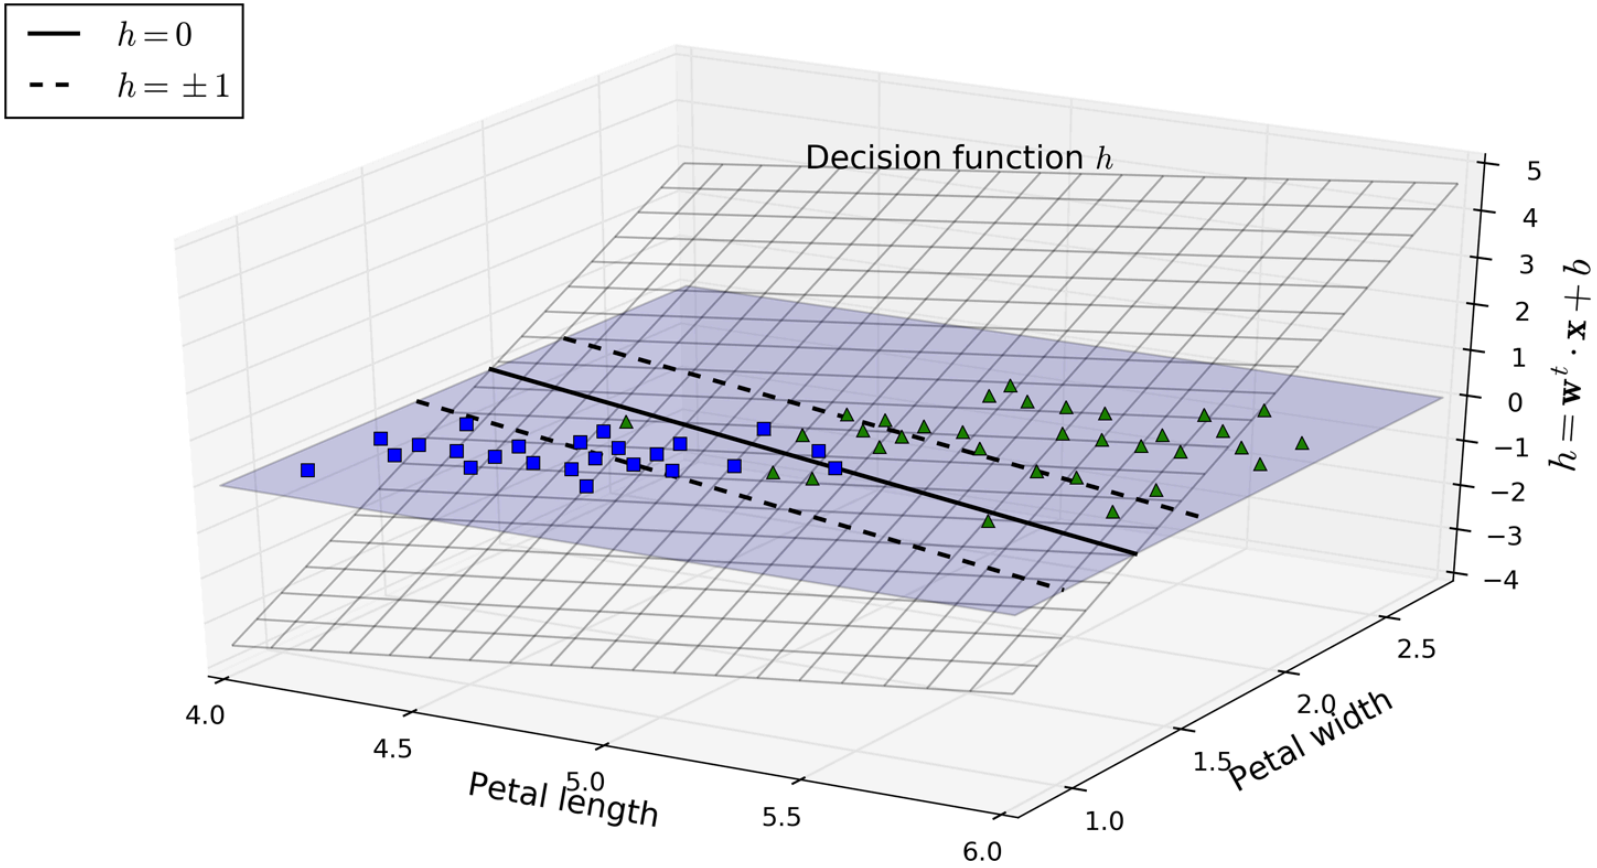

Linear Classification (cont'd)

- If the data is linearly separable, the SVM algorithm finds w and b, so that:

  Class +1 satisfies $w^T \cdot x + b \geq 1$

  Class -1 satisfies $w^T \cdot x + b \leq -1$

- Once w, b are known, new points x are classified by:

$$y = \begin{cases} +1, & if\ w^T \cdot x + b > 0 \\ -1, & if\ w^T \cdot x + b < 0 \end{cases}$$

Valid SVM solution

- What does a valid SVM solution look like?

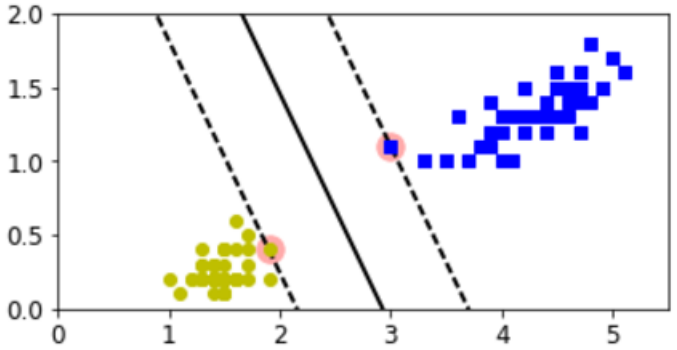

- What does it *not* look like?

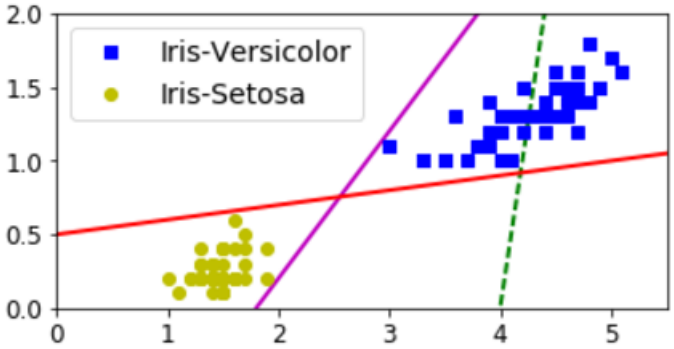

Valid SVM solution cont'd

- What does an SVM solution look like?

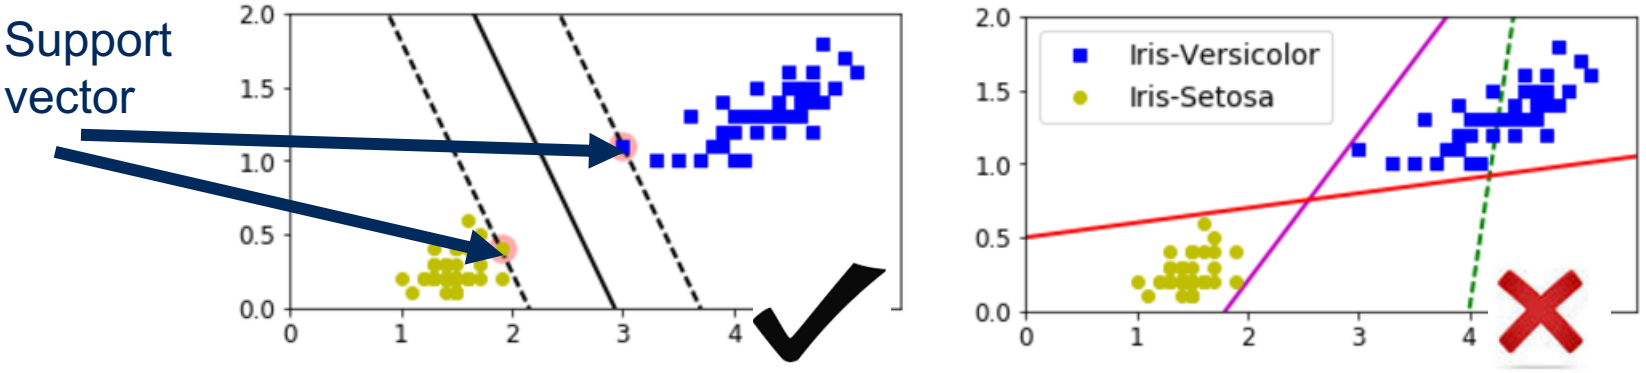

- The vectors on the dashed lines are the *support vectors,* and determine the solution
- Neither Setosa examples to the far left, or Versicolor examples to the far right, determine the solution

Hard Margin SVM

- SVM described below is a ***hard margin***
- No examples between dashed lines

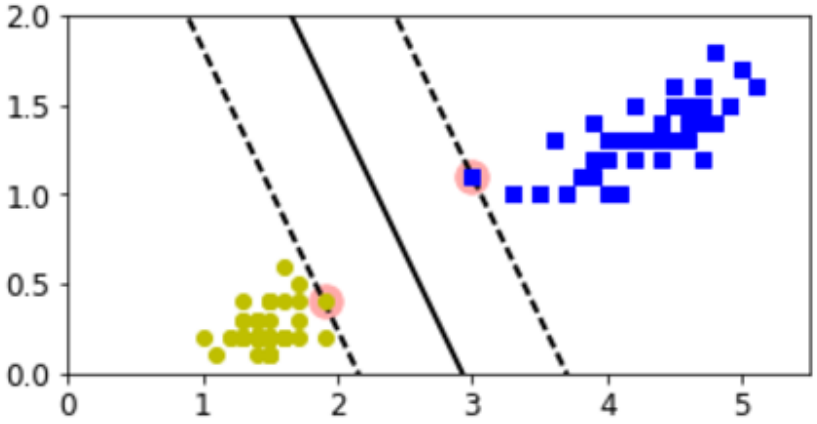

Soft Margins SVM

- ***Soft margin*** SVM allows violations
- Violation means:
  - having examples within the dashed lines
  - or even examples on the wrong side of the solid line
- The solution finds a balance between wide margin and number of violations

Triangle:
Violation and
Misclassification

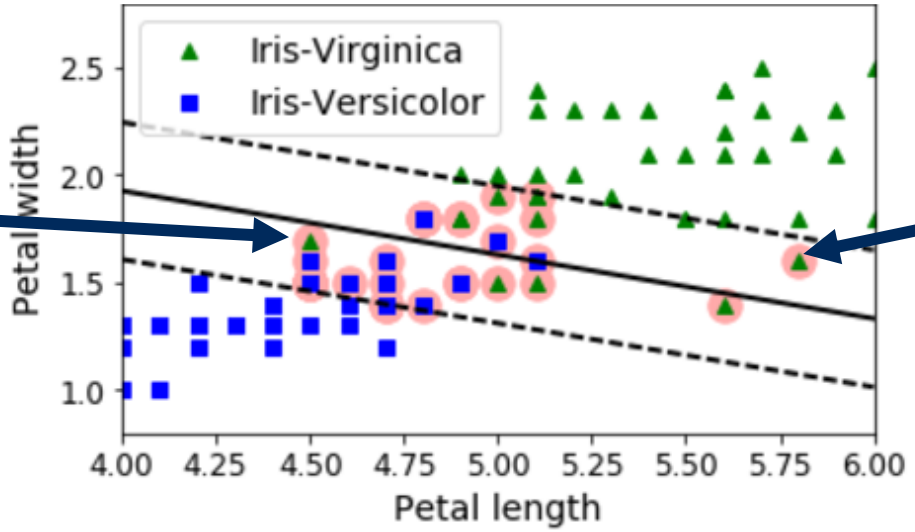

Triangle:
Violation but not
Misclassification

Controlling Margins in SVM

- Controlling margin is managed by a hyperparameter C
- C is the cost of margin violation
- Margin violation is bad
- Low C ⇔ big margin & more violations
- Left model C = 1 and right one is C = 100

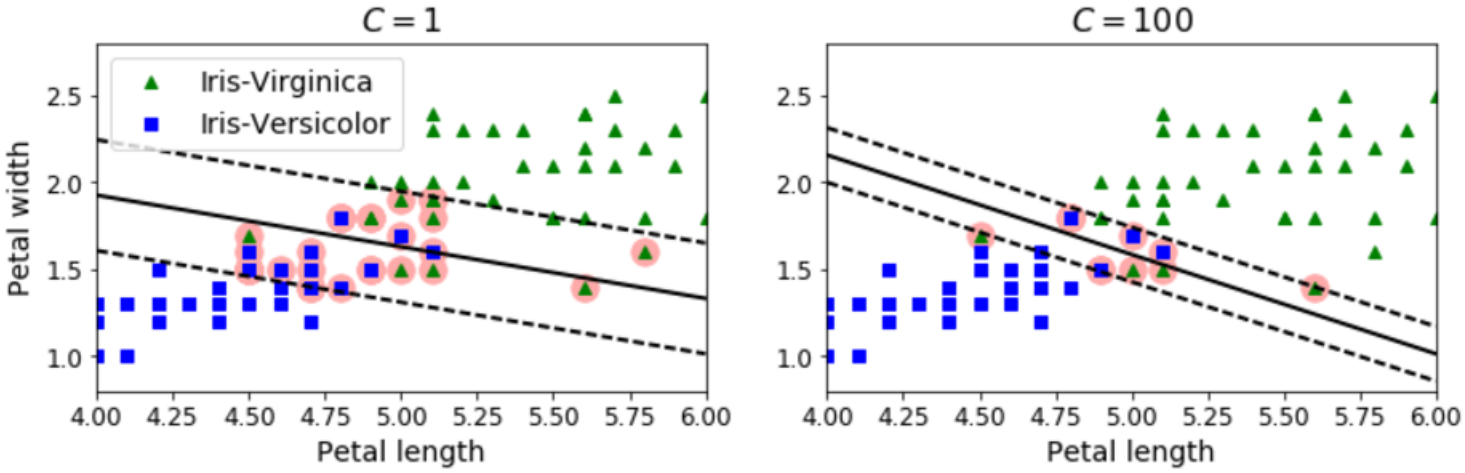

Overfitting and regularization

- High C (cost) leads to overfitting
- C can be used for regularization
- If SVM overfits, decrease C. It allows violations, less aggressively fitting the data
- Violations may not be misclassifications

Feature Scaling

- SVM is sensitive to feature scaling

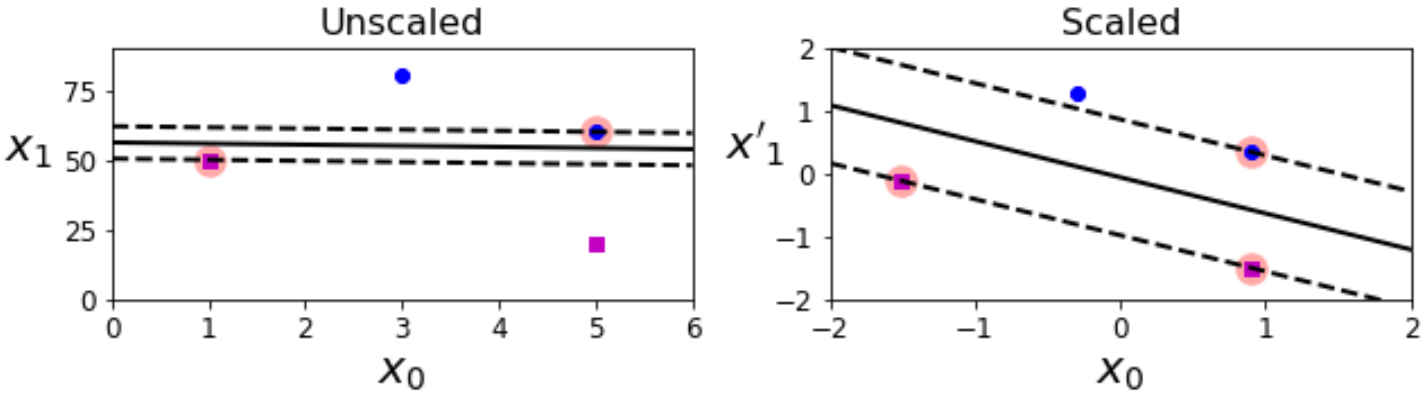

- Highly recommended to scale features for better solution

Outliers and SVM

- Hard-margin SVM is sensitive to outliers

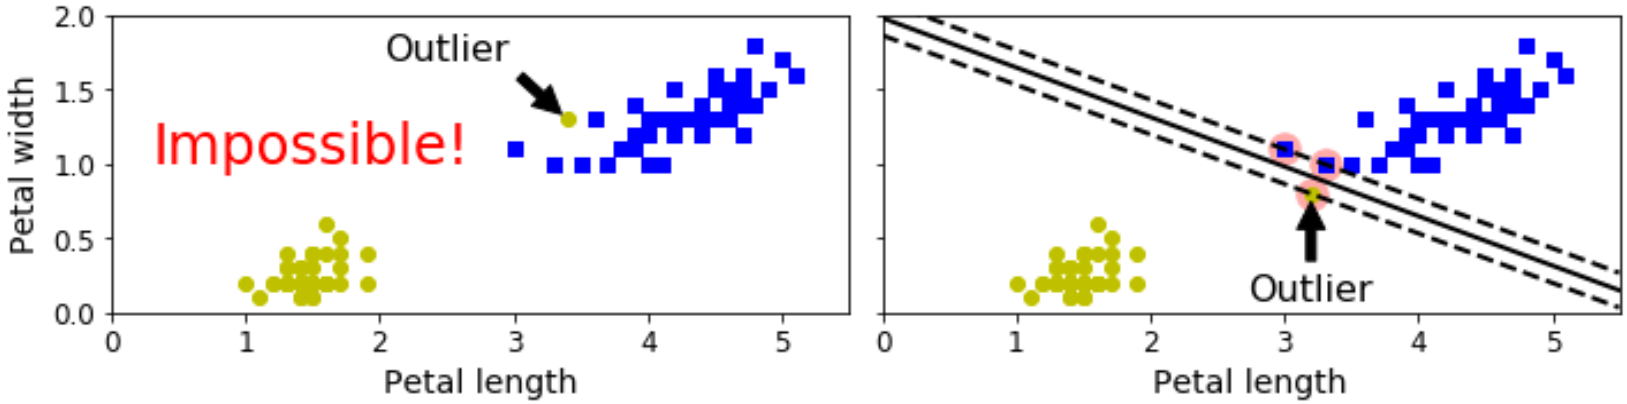

- Highly recommended to identify and remove outliers before training if you pick hard-margin (no violation allowed)

Module 6 – Section 3

Non-Linear Classification

Non-Linear Classification

- Non linearly-separable datasets can become so by adding new features
- A non linearly separable 1-feature set can become separable by transforming the feature (in this example, by squaring it)

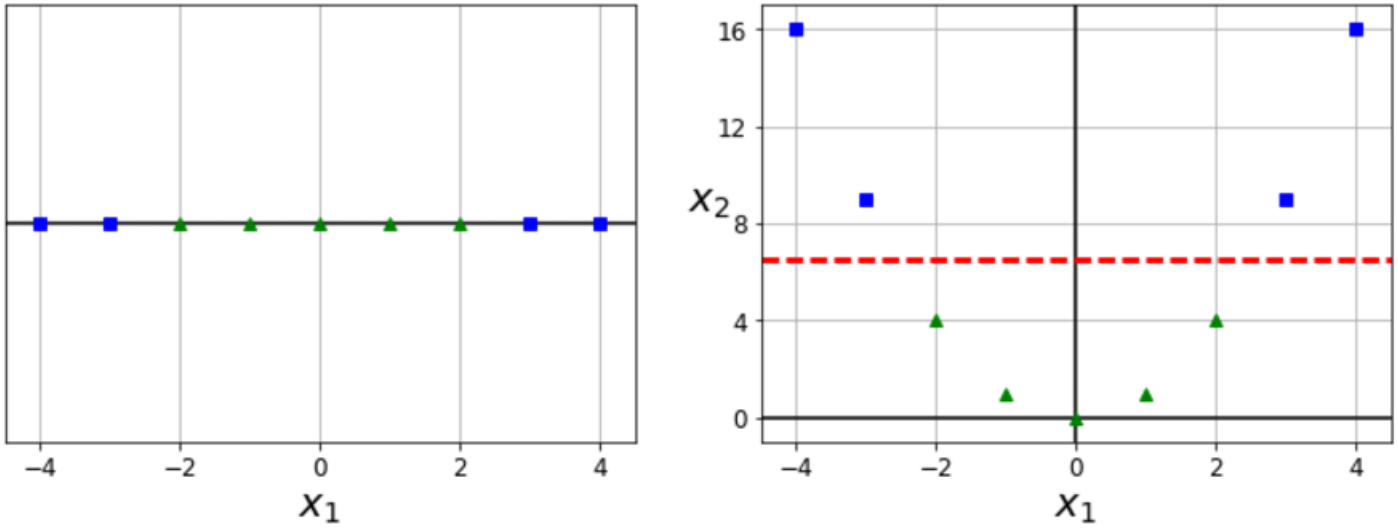

Non-Linear Classification (cont'd)

- Adding non-linear features manually is not the best approach.
- For large features, adding higher degree polynomial terms is time consuming, complicates model and slows down training
- There is a better way !

Non-Linear Classification (cont'd)

- The algorithm works as if the original features are transformed non-linearly into a high dimensional space

$$\phi: R^n \rightarrow R^r, r \gg n$$

- In R^r the dataset (hopefully) becomes separable, and linear SVM is applied
- The transformation ϕ is unknown. The algorithm does not actually use it
- In order to perform transformation without knowing ϕ, we use the Kernel trick

Kernel Trick

Kernel trick:

- Kernel is a magic mathematical method
- Without adding any polynomial term, kernel trick produces the same results as if you had manually added many polynomial terms
- So, no overfitting and simpler model
- Only need a kernel in the original space $k: R^n \times R^n \rightarrow R$. Mercer's Theorem. assures there exists $\phi: R^n \rightarrow R^r$ (for some r) such that $k(x, x') = \langle \phi(x), \phi(x') \rangle$
- To do linear SVM is the transformed space we only need inner product $\langle \phi(x), \phi(x') \rangle$

Kernel Trick (cont'd)

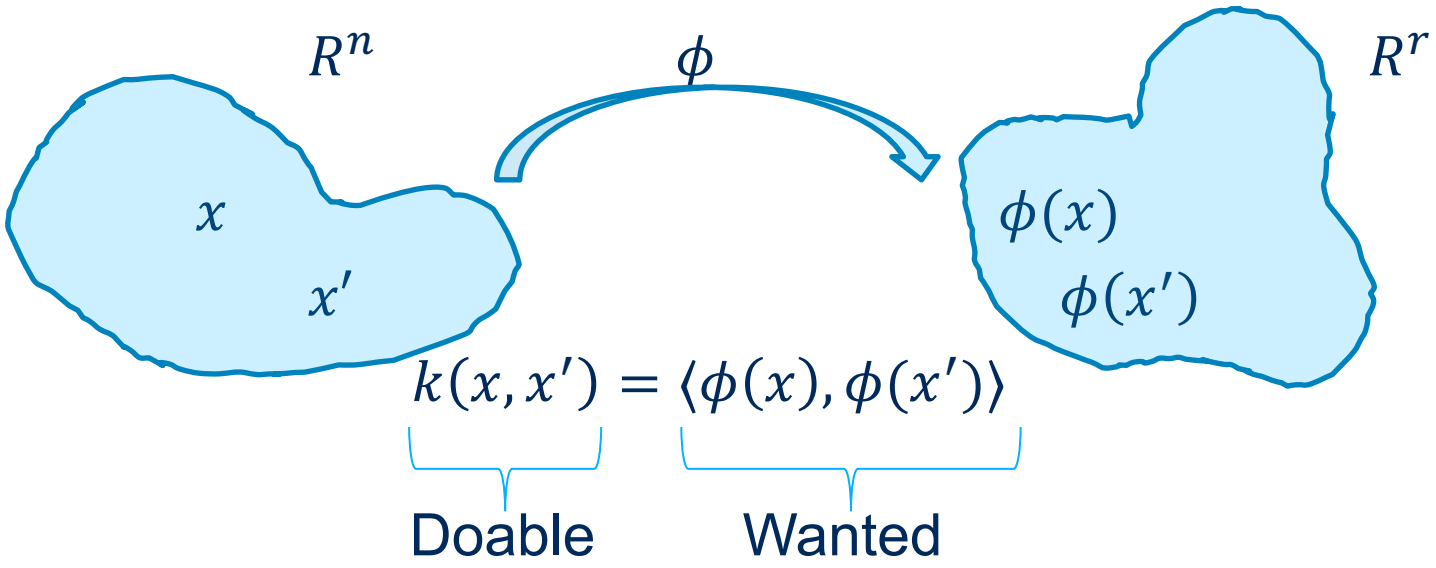

- Non-linear SVM requires specifying the kernel
- There are many possible kernels (hyperparameter)
- Popular kernels: polynomial, gaussian, linear

Dual problem

- In order to use Kernel trick, we need dual problem
- Given a constrained optimization problem (primal problem), it is possible to to express a different but closely related problem called its dual problem
- Under some conditions, the solution to dual problem can be the same as primal problem
- SVM meets those conditions and hence we can solve dual problem instead of the primal problem
- To better understand, imagine doing multiplication using summation operation

Available Kernels

- Linear: $k(x, x') = \langle x, x' \rangle$
- Polynomial: $k(x, x') = (\langle x, x' \rangle + r)^d$
- Gaussian RBF: $k(x, x') = \exp\left(-\gamma \left|\left|x - x'\right|\right|^2\right) = \exp(-\gamma \langle x - x', x - x' \rangle)$
- The kernel k, and its parameters, e.g. d, γ, etc. need to be calibrated throughout cross-validation (e.g., train-valid, or k-fold)

Polynomial Kernel

- Polynomial kernel is analogous to replacing the original features by many polynomial transformations of them
- d is the desired polynomial dimension

$$k(x, x') = (\langle x, x' \rangle + r)^d$$

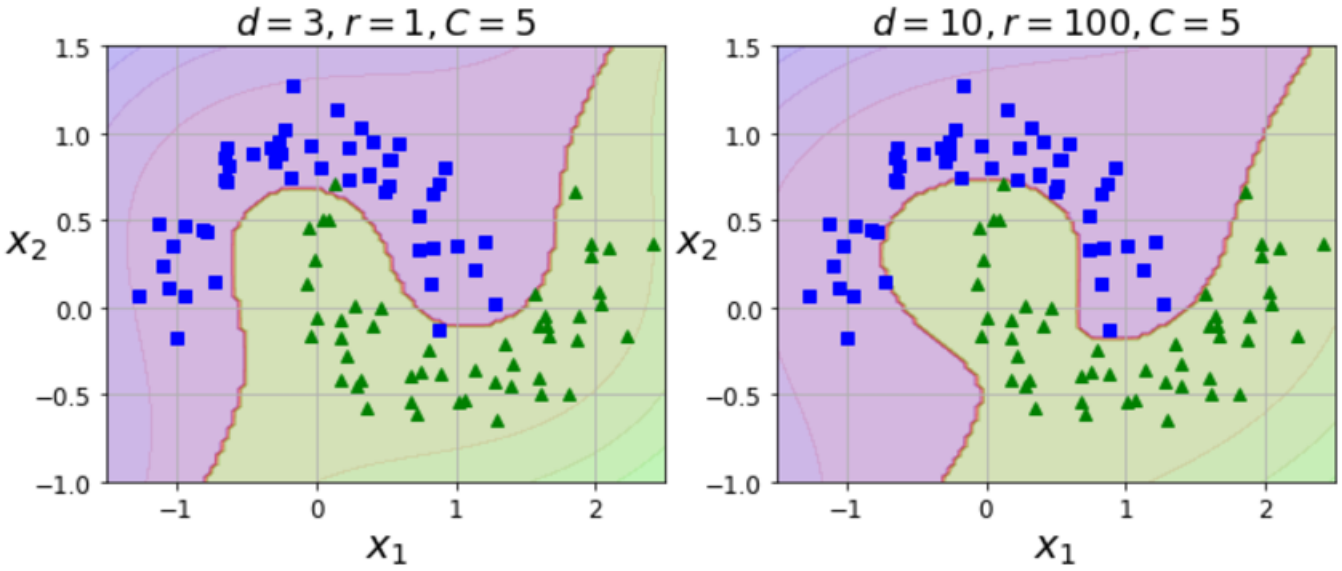

Gaussian RBF Kernel

- Gaussian RBF Kernel

$$k(x, x') = \exp(-\gamma \left|\left|x - x'\right|\right|^2)$$

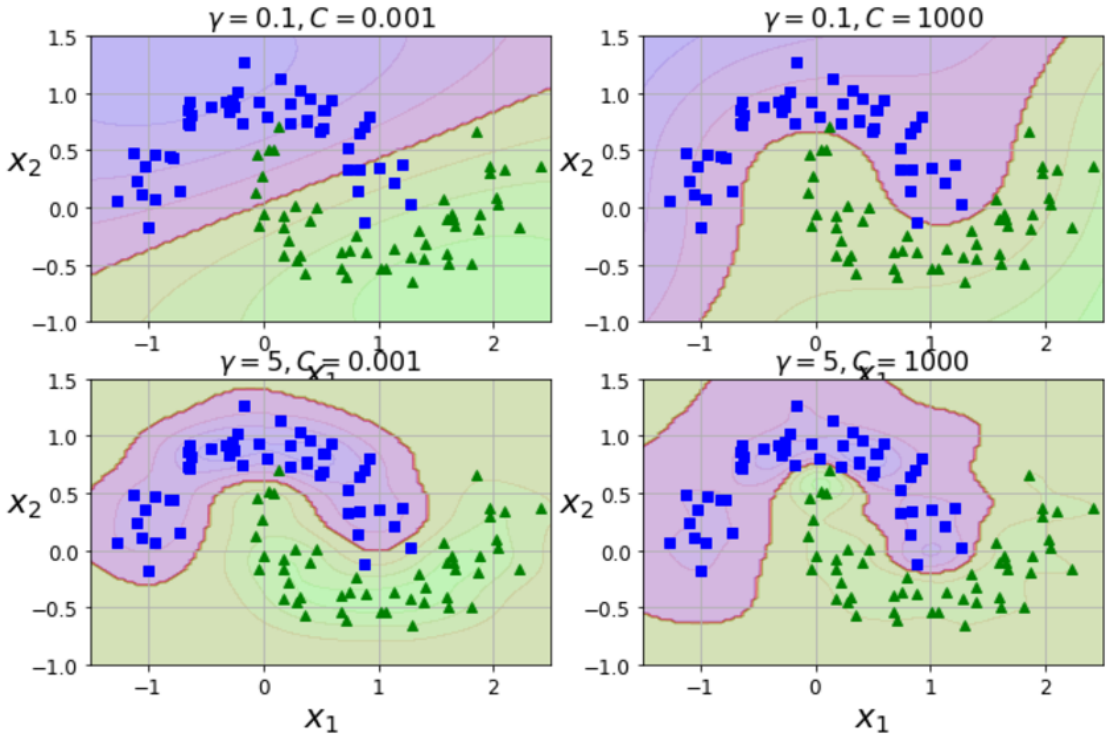

Module 6 – Section 4

SVM Regression

SVM Regression

- Classification goal: find parallel dashed lines as far as possible and containing few points between them
- Regression goal: find parallel dashed lines that are close to each other and contain as many points as possible
- Instead of C the regularization hyperparameter is ϵ

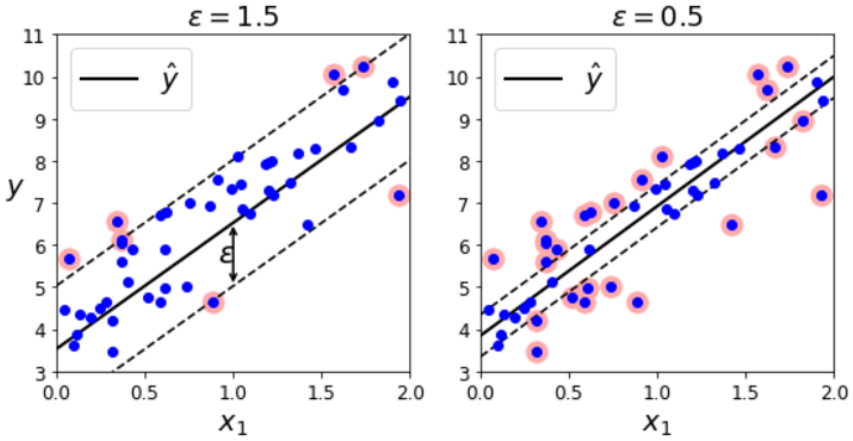

SVM Regression (cont'd)

- Non linear SVM Regression, same kernels
- E.g., polynomial kernels:

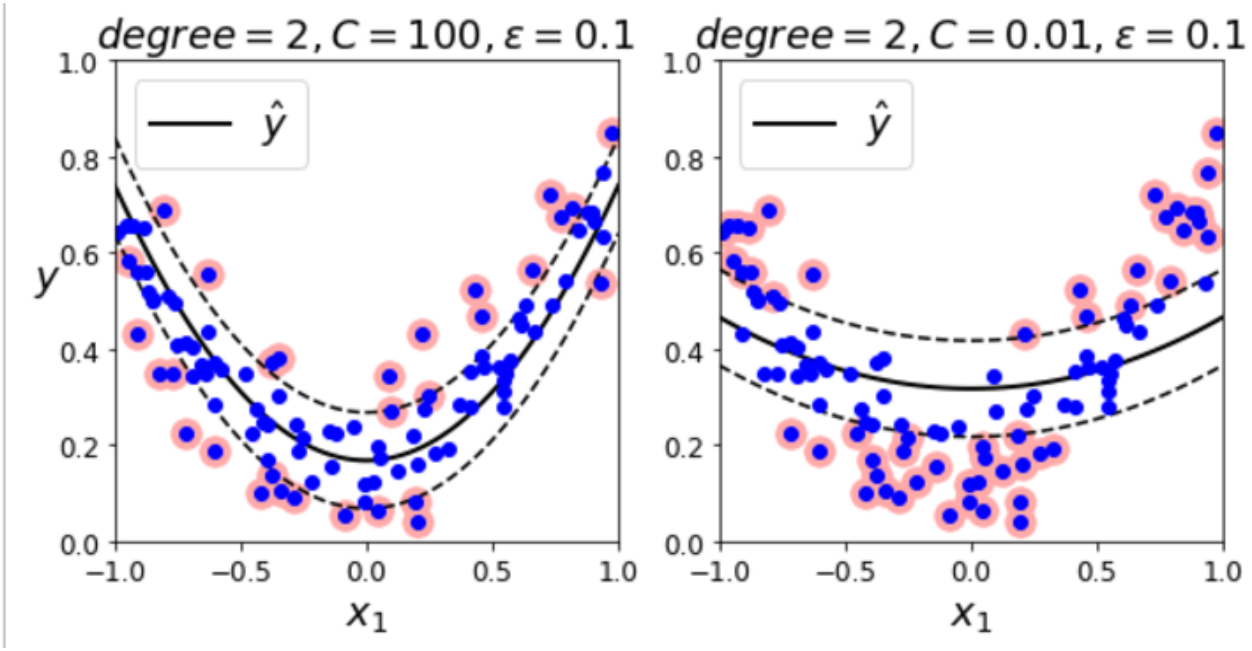

Module 6 – Section 5

How SVM Works

How SVM works

- Linear SVM classifies points by:

$$y = \begin{cases} +1, & if\ w^T \cdot x + b > 0 \\ -1, & if\ w^T \cdot x + b < 0 \end{cases}$$

- The algorithm finds w ($\perp\ to\ line$) and b

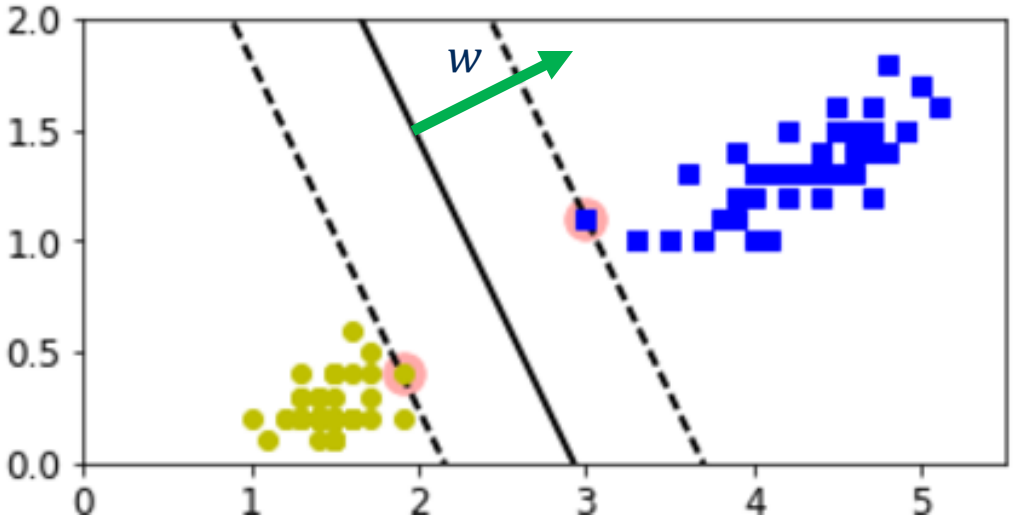

How SVM works (cont'd)

- Linear SVM classifies points by:

$$y = \begin{cases} +1, & if\ w^T \cdot x + b > 0 \\ -1, & if\ w^T \cdot x + b < 0 \end{cases}$$

- The algorithm finds w ($\perp\ to\ line$) and b

- Distance between margins: $\frac{w}{||w||}(x^+ - x^-) = \frac{2}{||w||}$

x^+, x^- are support vectors

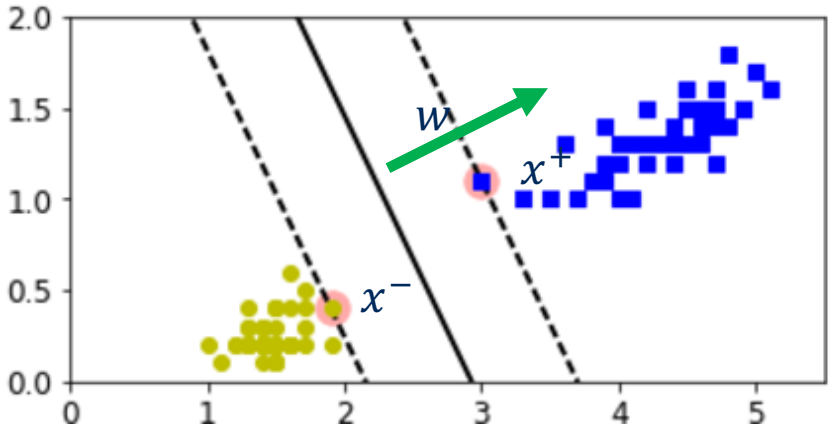

How SVM works (cont'd)

- Optimization problem that yields w, b

$$\begin{cases} minimize: \frac{1}{2} w^T \cdot w \\ subject\ to: \ y^{(i)}\left(w^T \cdot x^{(i)} + b\right) \geq 1, for\ all\ i \end{cases}$$

- For soft margin: add slack variables

$$\begin{cases} minimize: \frac{1}{2} w^T \cdot w + C \sum_i \xi^{(i)} \\ subject\ to: y^{(i)}\left(w^T \cdot x^{(i)} + b\right) \geq 1 - \xi^{(i)} \end{cases}$$

How SVM works (cont'd)

- These formulations (*Primal*) are cases of Quadratic Programming

$$\begin{cases} minimize\ (variable\ u)\colon \dfrac{1}{2}u^T \cdot H \cdot u + f^T \cdot u \\ subject\ to\colon A \cdot u \leq d \end{cases}$$

where: $u, f \in R^{q\times 1}, H\ symmetrix\ q\times q, d \in R^{m\times 1}, A \in R^{m\times q}$

- With appropriate values for H, f, A, d, the primal formulation can be solved as QP problems (well studies)

How SVM works (cont'd)

- There is an equivalent formulation (same solution) to the Primal, called *Dual*

$$\begin{cases} minimize\ (variable\ \alpha): \sum_{i=1}^{m} \alpha_i y^{(i)} - \frac{1}{2} \sum_{i,j=1}^{m} \alpha_i \alpha_j y^{(i)} y^{(j)} \langle x^{(i)}, x^{(j)} \rangle \\ subject\ to: \alpha_i \geq 0, \sum_{i=1}^{m} \alpha_i y^{(i)} = 0 \end{cases}$$

- Once α is found, it can be used to obtain w and b from the Primal formulation, and build the classifier

$$b = \frac{1}{n_s} \sum_{i:\alpha_i > 0} (1 - y^{(i)} \langle w, x^{(i)} \rangle)$$

$$w = \sum_{i=1}^{m} \alpha_i y^{(i)} x^{(i)}$$

How SVM works (cont'd)

- The Dual formulation is also applicable for the non-linear case, using any kernel: replace $\langle \_, \_ \rangle$ by k

$$\begin{cases} minimize\ (variable\ \alpha): \sum_{i=1}^{m} \alpha_i y^{(i)} - \frac{1}{2} \sum_{i,j=1}^{m} \alpha_i \alpha_j y^{(i)} y^{(j)} (k(x^{(i)}, x^{(j)})) \\ subject\ to: \alpha_i \geq 0, \sum_{i=1}^{m} \alpha_i y^{(i)} = 0 \end{cases}$$

- Once α is found, b can be found, and the classifier can be computed $y = w^T \cdot \phi(x) + b$:

$$b = \frac{1}{n_s} \sum_{i:\alpha_i > 0} \left(1 - y^{(i)} \sum_{j=1}^{m} \alpha_j y^{(j)} k(x^{(j)}, x^{(i)}) \right)$$

$$y = w^T x + b = \sum_{i=1}^{m} \alpha_i y^{(i)} k(x^{(i)}, x) + b$$

UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 6 – Section 6

Resources and Wrap-up

Resources

- Hands-On Machine Learning with Scikit-Learn and Tensorflow

Assessment

- See Jupyter Notebook

Next Class

- Decision Trees

Follow us on social

Join the conversation with us online:

facebook.com/uoftscs

@uoftscs

linkedin.com/company/university-of-toronto-school-of-continuing-studies

@uoftscs

Any questions?

UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Thank You

Thank you for choosing the University of Toronto School of Continuing Studies